

KEY INSERTION SEQUENCE: B C G E F D A ~~F~~ J.

# SPLITTING AND PROMOTING



FIGURE 9.14 ■ Initial leaf of a B-tree with a page size of seven.



FIGURE 9.15 ■ Splitting the leaf to accommodate the new *J* key.

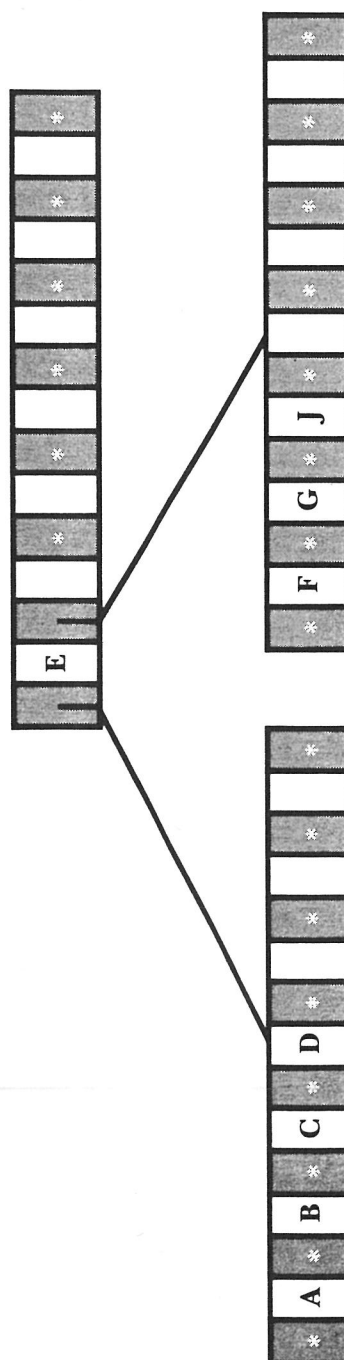


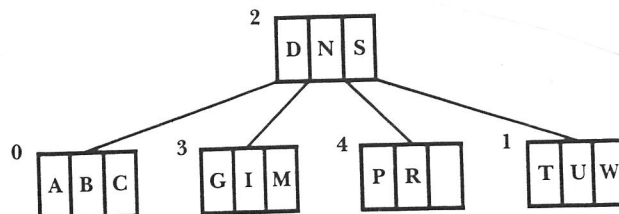
FIGURE 9.16 ■ Promotion of the *E* key into a root node.

## ALGORITHMS FOR B-TREE SEARCHING

```
FUNCTION: search (RRN, KEY, FOUND_RRN , FOUND_POS)

  if RRN == NIL then    /* stopping condition for the recursion */
    return NOT FOUND
  else
    read page RRN into PAGE
    look through PAGE for KEY, setting POS equal to the
      position where KEY occurs or should occur.
    if KEY was found then
      FOUND_RRN := RRN      /* current RRN contains the key */
      FOUND_POS := POS
      return FOUND
    else /* follow CHILD reference to next level down */
      return(search(PAGE.CHILD[POS], KEY, FOUND_RRN, FOUND_POS))
    endif
  endif
end FUNCTION
```

**FIGURE 9.20** • Function *search (RRN, KEY, FOUND\_RRN, FOUND\_POS)* searches recursively through the B-tree to find KEY. Each invocation searches the page referenced by RRN. The arguments FOUND\_RRN and FOUND\_POS identify the page and position of the key, if it is found. If *search()* finds the key, it returns FOUND. If it goes beyond the leaf level without finding the key, it returns NOT FOUND.



**FIGURE 9.21** • B-tree used for the search example.

```

FUNCTION: insert (CURRENT_RRN, KEY PROMO_R_CHILD, PROMO_KEY)

  if CURRENT_RRN = NIL then      /* past bottom of tree */
    PROMO_KEY := KEY
    PROMO_R_CHILD := NIL
    return PROMOTION
  else
    read page at CURRENT_RRN into PAGE
    search for KEY in PAGE.
    let POS := the position where KEY occurs or should occur.

    if KEY found then
      issue error message indicating duplicate key
      return ERROR

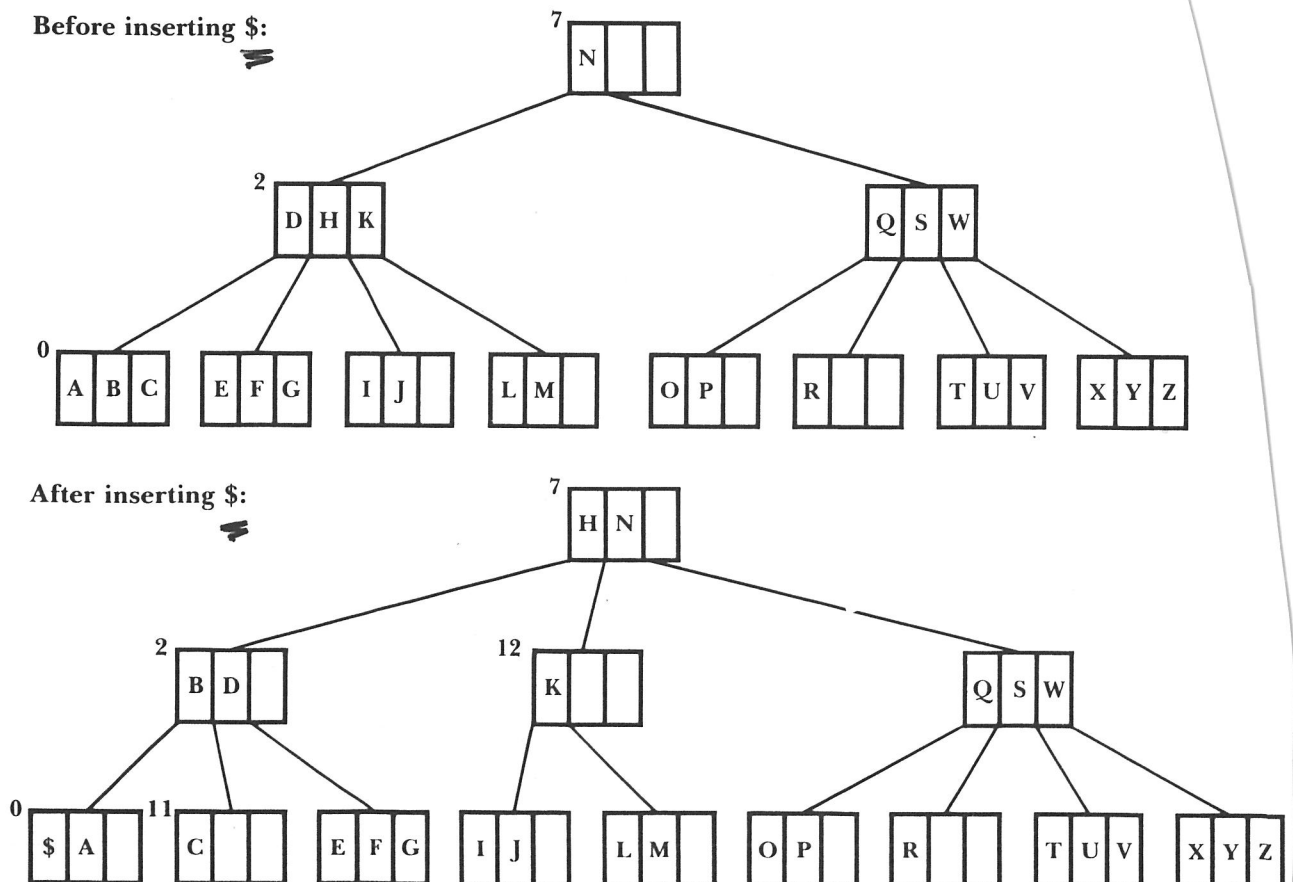
    RETURN_VALUE := insert(PAGE.CHILD[POS], KEY, P_B_RRN, P_B_KEY)

    if RETURN_VALUE == NO PROMOTION or ERROR then
      return RETURN_VALUE

    elseif there is space in PAGE for P_B_KEY then
      insert P_B_KEY and P_B_RRN (promoted from below) in PAGE
      return NO PROMOTION
    else
      split(P_B_KEY, P_B_RRN, PAGE, PROMO_KEY, PROMO_R_CHILD, NEWPAGE)
      write PAGE to file at CURRENT_RRN
      write NEWPAGE to file at rrn PROMO_R_CHILD
      return PROMOTION /* promoting PROMO_KEY and PROMO_R_CHILD */
    endif
  endif
end FUNCTION

```

**FIGURE 9.24** • Function *insert* (*CURRENT\_RRN*, *KEY*, *PROMO\_R\_CHILD*, *PROMO\_KEY*) inserts a *KEY* in a B-tree. The insertion attempt starts at the page with relative record number *CURRENT\_RRN*. If this page is not a leaf page, the function calls itself recursively until it finds *KEY* in a page or reaches a leaf. If it finds *KEY*, it issues an error message and quits, returning *ERROR*. If there is space for *KEY* in *PAGE*, *KEY* is inserted. Otherwise, *PAGE* is split. A split assigns the value of the middle key to *PROMO\_KEY* and the relative record number of the newly created page to *PROMO\_R\_CHILD* so that insertion can continue on the recursive ascent back up the tree. If a promotion does occur, *insert()* indicates this by returning *PROMOTION*. Otherwise, it returns *NO PROMOTION*.



**FIGURE 9.22** • The effect of adding \$ to the tree constructed in Fig. 9.18.

PROCEDURE: split (I\_KEY, I\_RRN, PAGE, PROMO\_KEY, PROMO\_R\_CHILD, NEWPAGE)

copy all keys and pointers from PAGE into a working page that can hold one extra key and child.

insert I\_KEY and I\_RRN into their proper places in the working page.

allocate and initialize a new page in the B-tree file to hold NEWPAGE.

set PROMO\_KEY to value of middle key, which will be promoted after the split.

set PROMO\_R\_CHILD to RRN of NEWPAGE.

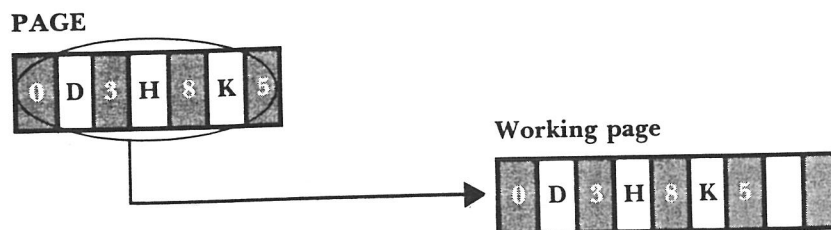
copy keys and child pointers *preceding* PROMO\_KEY from the working page to PAGE.

copy keys and child pointers *following* PROMO\_KEY from the working page to NEWPAGE.

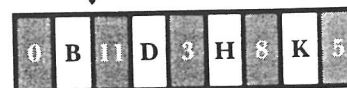
end PROCEDURE

**FIGURE 9.25 • Split (I\_KEY, I\_RRN, PAGE, PROMO\_KEY, PROMO\_R\_CHILD, NEWPAGE)**, a procedure that inserts I\_KEY and I\_RRN, causing overflow, creates a new page called NEWPAGE, distributes the keys between the original PAGE and NEWPAGE, and determines which key and RRN to promote. The promoted key and RRN are returned via the arguments PROMO\_KEY and PROMO\_R\_CHILD.

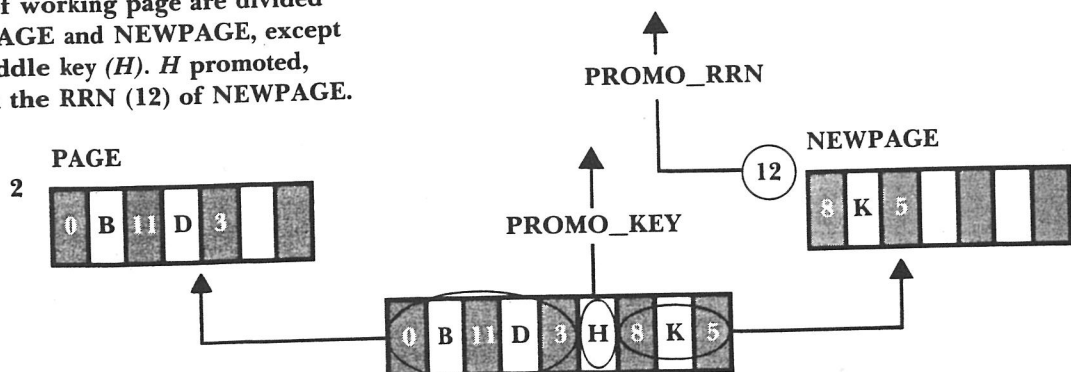
Contents of PAGE are copied to the working page.



I\_KEY (B) and I\_RRN (11) are inserted into working page.



Contents of working page are divided between PAGE and NEWPAGE, except for the middle key (H). H promoted, along with the RRN (12) of NEWPAGE.

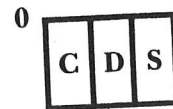


**FIGURE 9.26 • The movement of data in split().**

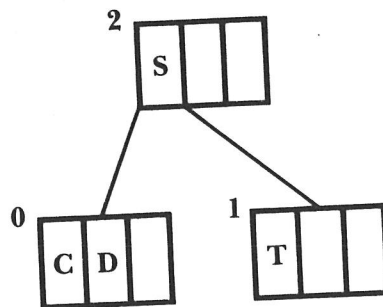
# KEY INSERTION SEQUENCE:

C S D T  
A M P I  
B W N G  
U R K E  
H O L J  
Y Q Z F  
X V

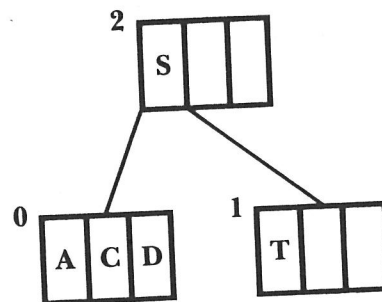
Insertion of C, S, and D  
into the initial page:



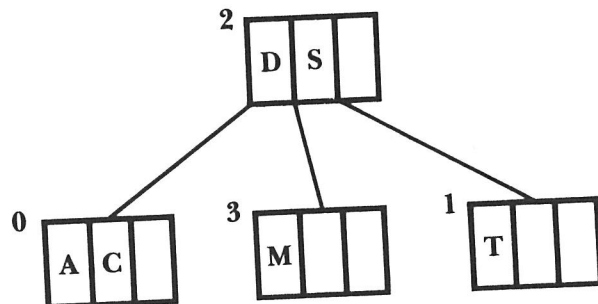
Insertion of T forces the split  
and the promotion of S:



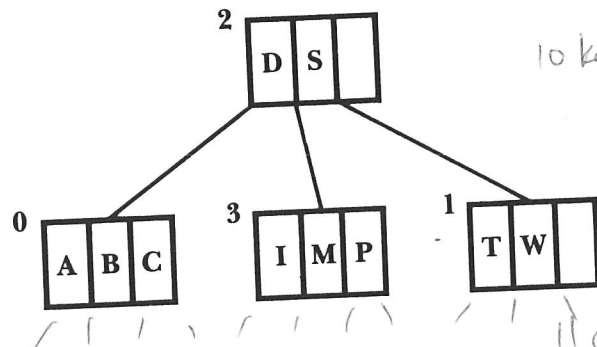
A added without incident:



Insertion of M forces another  
split and the promotion of D:



P, I, B, and W inserted  
into existing pages:



Insertion of N causes another  
split, followed by the promo-  
tion of N. G, U, and R are  
added to existing pages:

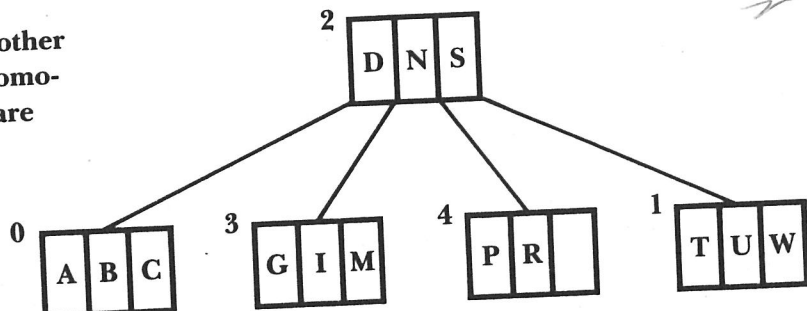
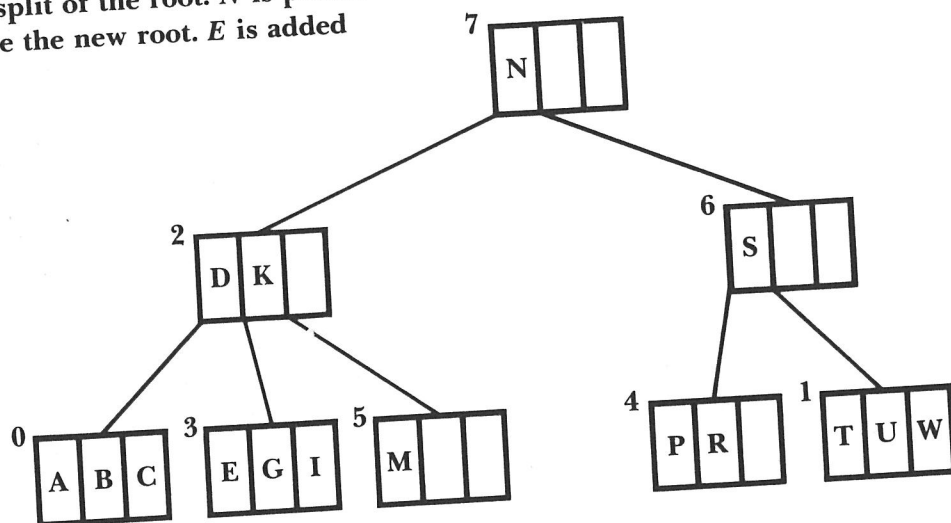
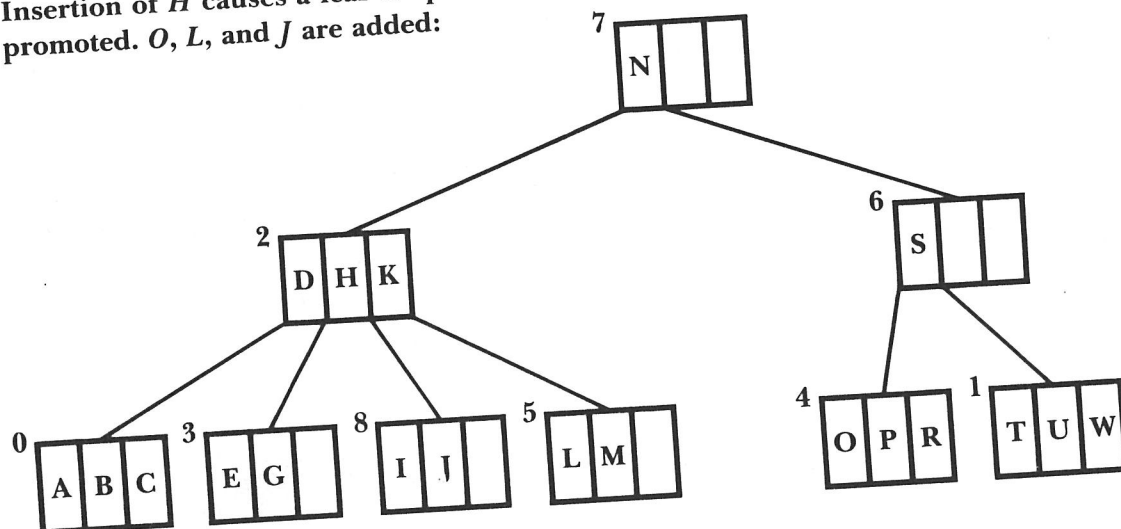


FIGURE 9.17 • Growth of a B-tree, Part I. The tree grows to a point at which splitting of the root is imminent.

Insertion of *K* causes a split at leaf level, followed by the promotion of *K*. This causes a split of the root. *N* is promoted to become the new root. *E* is added to a leaf:



Insertion of *H* causes a leaf to split. *H* is promoted. *O*, *L*, and *J* are added:



Insertion of *Y* and *Q* force two more leaf splits and promotions. Remaining letters are added:

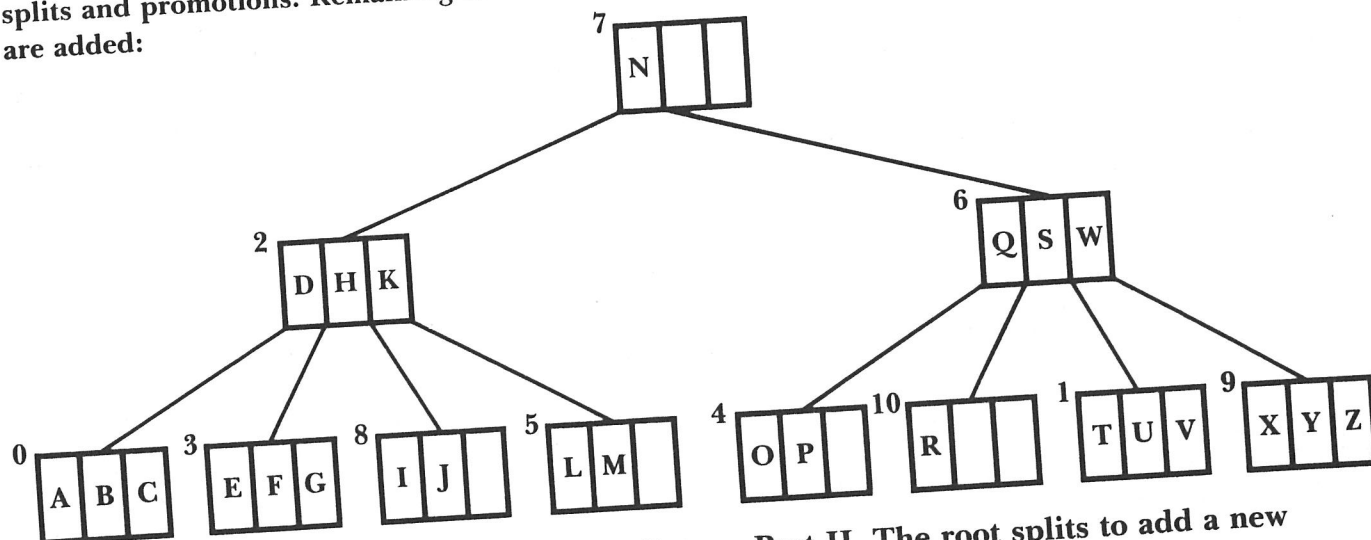
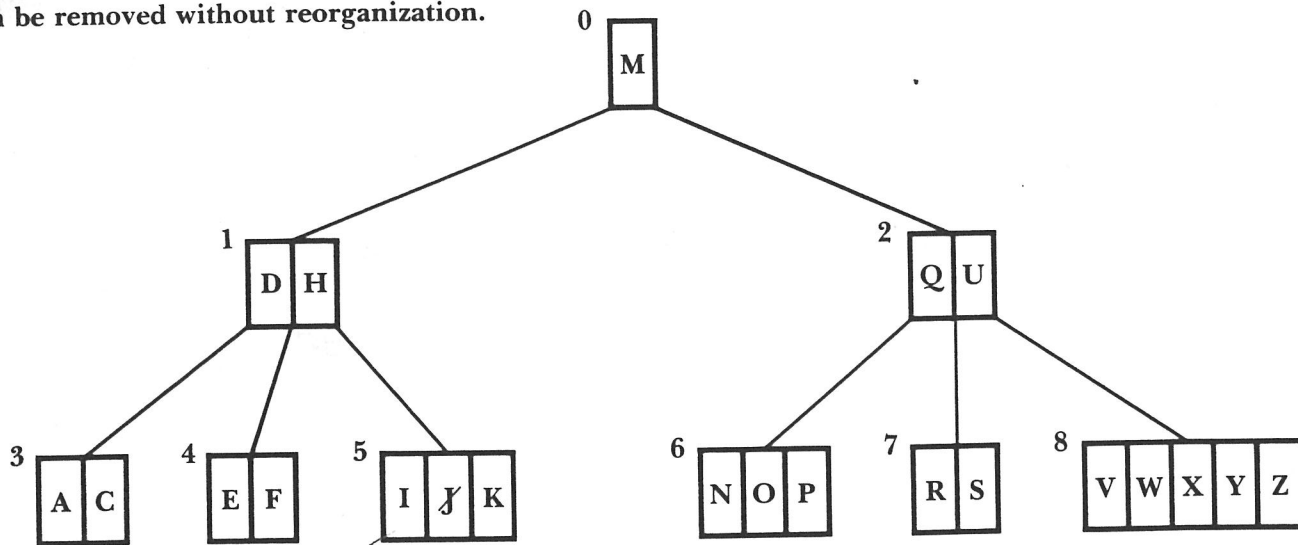


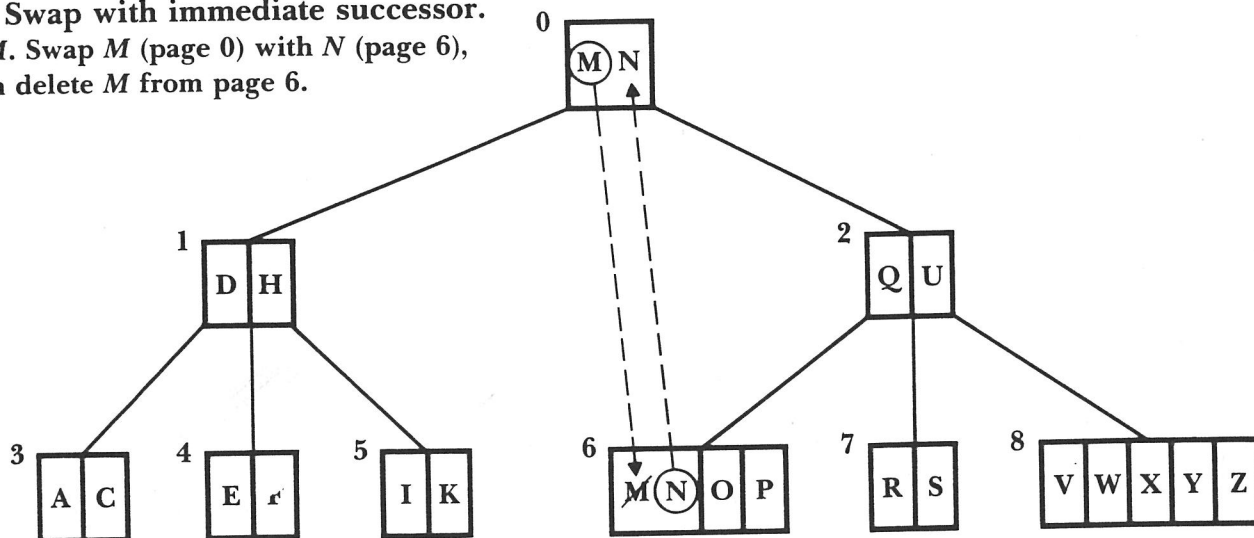
FIGURE 9.18 • Growth of a B-tree, Part II. The root splits to add a new level; remaining keys are inserted.

**Case 1: No action.**

Delete *J* from page 5. Since page 5 has more than the minimum number of keys, *J* can be removed without reorganization.



**Case 2: Swap with immediate successor.**  
Delete *M*. Swap *M* (page 0) with *N* (page 6), and then delete *M* from page 6.



**Case 3: Redistribution.**

Delete *R*. Underflow occurs. Redistribute keys among pages 2, 7, and 8 to restore balance between leaves.

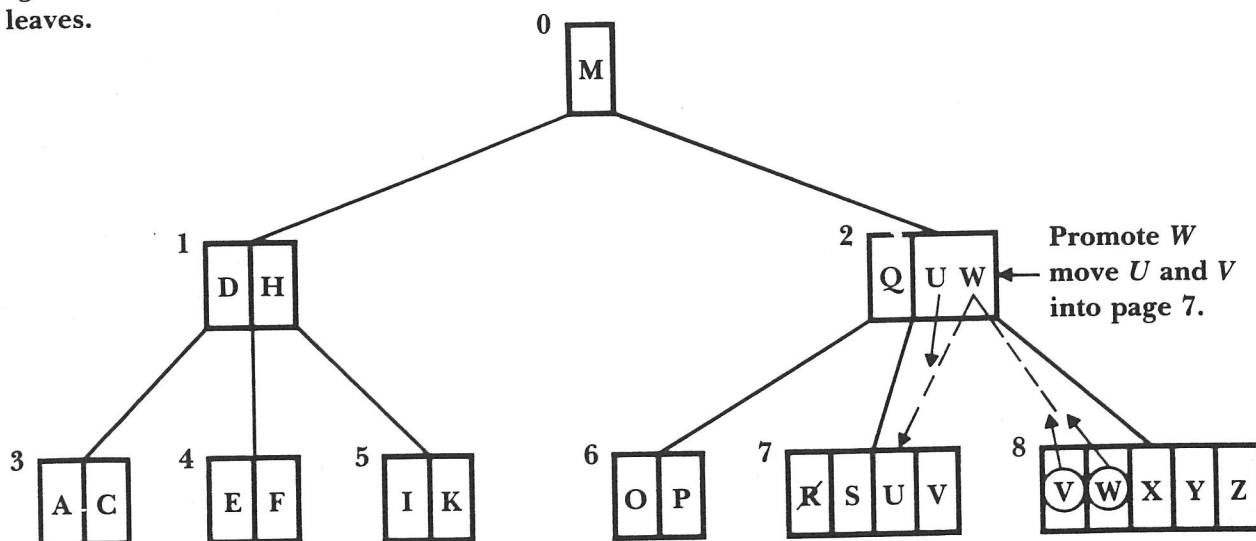
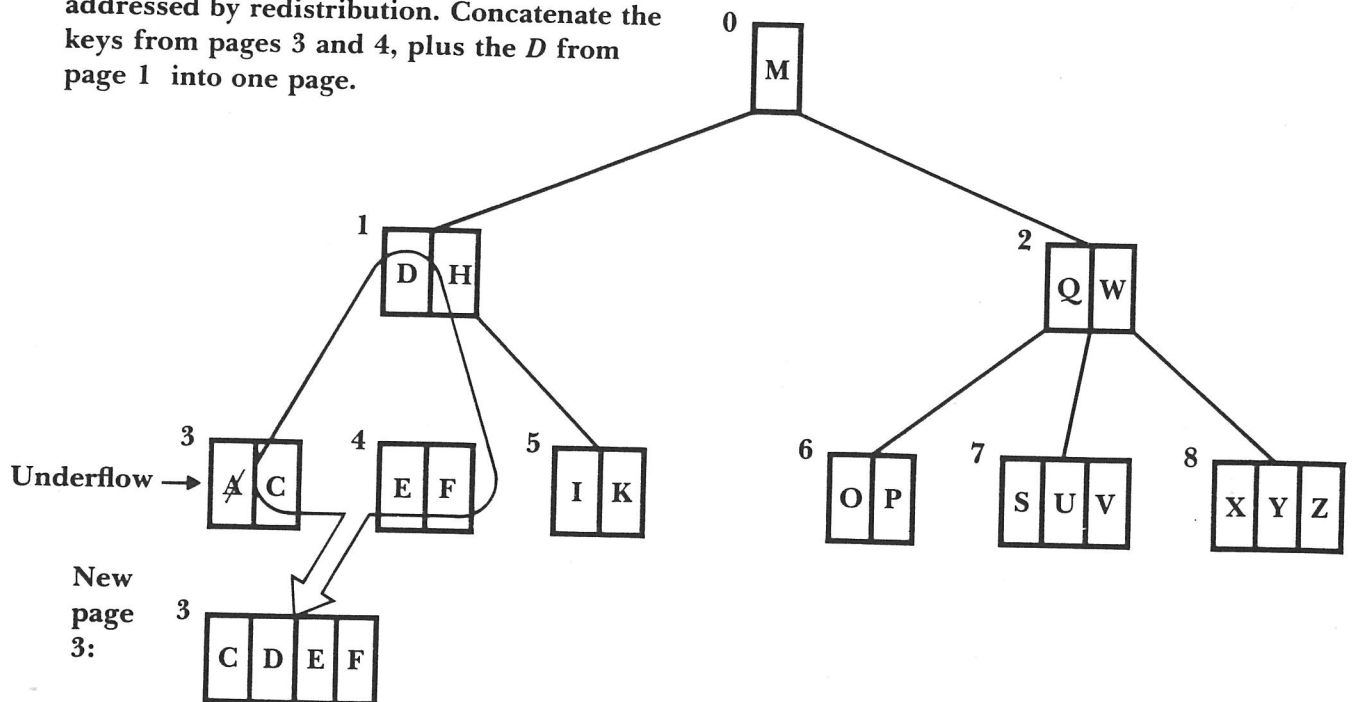


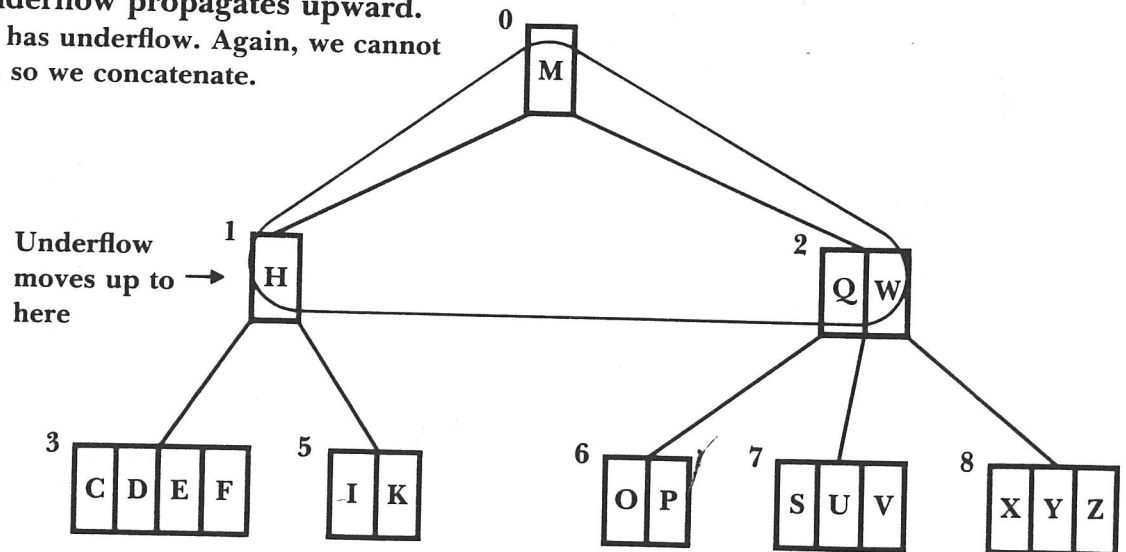
FIGURE 9.29 • Six situations that can occur during deletions.

#### Case 4: Concatenation.

Delete A. Underflow occurs, but it cannot be addressed by redistribution. Concatenate the keys from pages 3 and 4, plus the D from page 1 into one page.



Case 5: Underflow propagates upward. Now page 1 has underflow. Again, we cannot redistribute, so we concatenate.



#### Case 6: Height of tree decreased.

Since the root contains only one key, it is absorbed into the new root.

