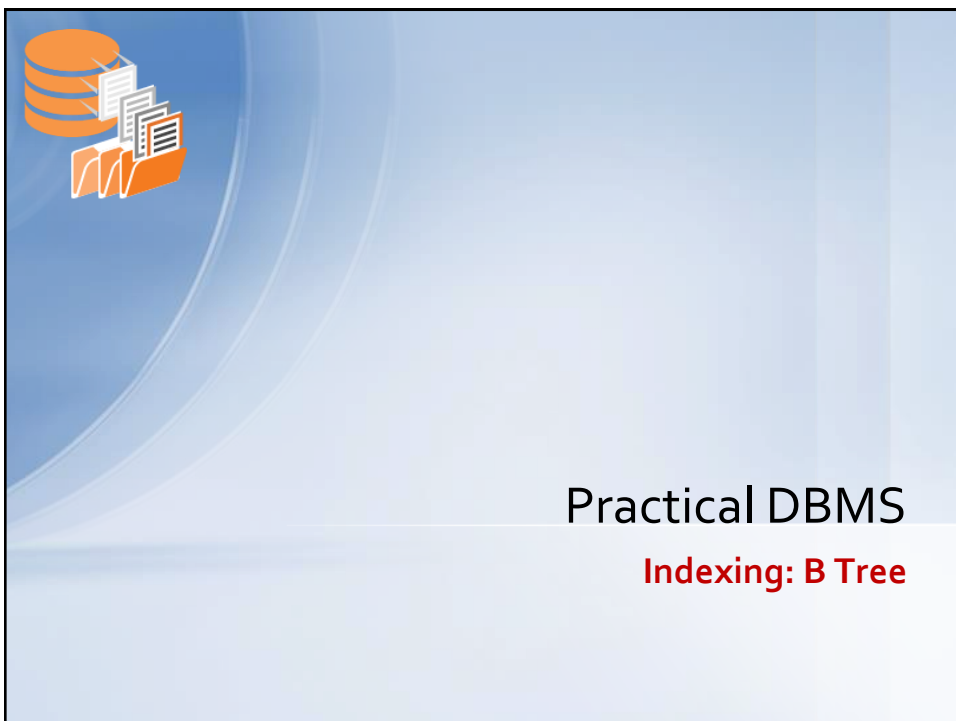



1

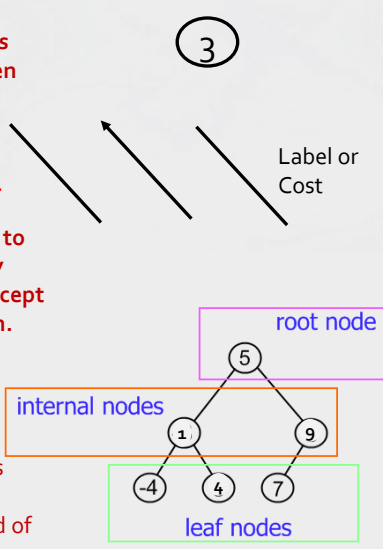


2



Tree Basics

- A **tree** is a data structure whose components are a set of **nodes** and a set of **edges** between the nodes.
 - Each node has value(s);
 - Each edge connects two nodes;
 - edges could be directed and labeled too.
- A **binary tree** has a directed edges (i.e. child to parent – 5 is parent of 3), each child has only one parent (except root), and each node (except leaf nodes) is a parent of one or two children.
- **Example:**
 - A root has no parent;
 - A leaf has no kids;
 - If node 5 is **level 0** then 3 is level 1, and 7 is level 2;
 - A **sub-tree** rooted at value -8 is composed of nodes -8, -4, and 9.




Joseph Vella - Practical DBMS

3

Tree Indexing - B+ Tree

3



Tree Searching (and Building): Effective only possible if placement is ordered

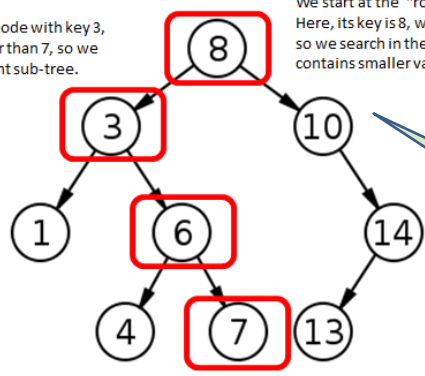
Insert / Find node with a 7 value.

Step 1:
We start at the “root” node. Here, its key is 8, which is larger than 7, so we search in the left sub-tree, which contains smaller values than 8.

Step 2:
We arrive at a node with key 3, which is smaller than 7, so we search in its right sub-tree.

Step 3:
We arrive at a node with key 6, which is again smaller than 7, so we again search in its right sub-tree.

Step 4:
By comparing each node’s key to the key we are looking for, we eventually arrive at the right node with key 7.




Joseph Vella - Practical DBMS

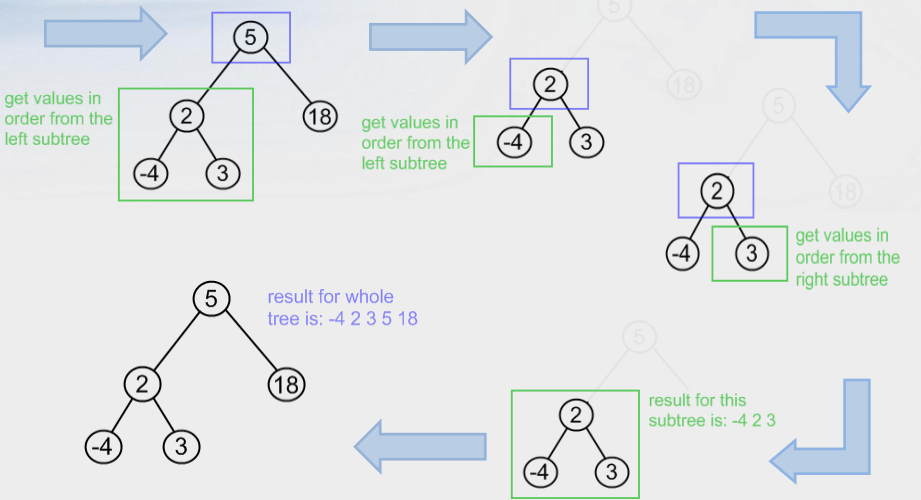
4

Tree Indexing - B+ Tree

4



Traversal over a Search Tree



get values in order from the left subtree

get values in order from the left subtree

get values in order from the right subtree

result for whole tree is: -4 2 3 5 18


result for this subtree is: -4 2 3

Joseph Vella - Practical DBMS

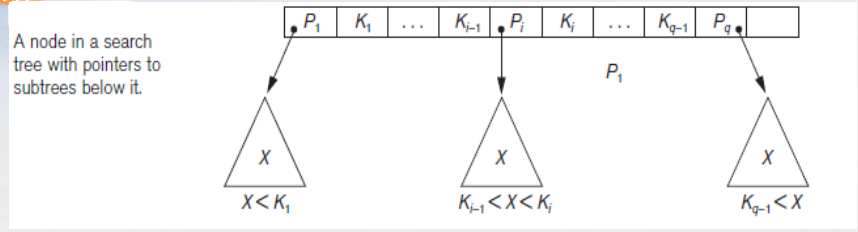
5

Tree Indexing - B+ Tree

5



Search Tree node structure



A node in a search tree with pointers to subtrees below it.

P_1

X
 $X < K_1$

X
 $K_{i-1} < X < K_i$

X
 $K_{g-1} < X$


- P_i are pointers and K_i are the keys;**
 - There are always **one** more pointer than keys;
 - Each key value represents an indexing key;
 - Each non-null pointer refers to a sub-tree of node structures;
- In case of a binary search tree, the structure has *two* pointers (i.e., P_1 and P_2) and one key value (i.e., K_1).**
- In the case of a B-tree the number of pointers and keys in a node is large!**
 - In which case the values of K_i in a page are ordered, i.e., $k_1 \leq k_2 \leq \dots \leq k_n$.
 - Note, that duplicate key values are not disallowed

Joseph Vella - Practical DBMS

6

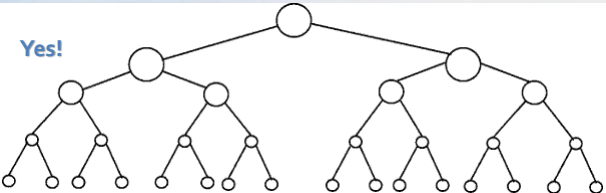
Tree Indexing - B+ Tree

6

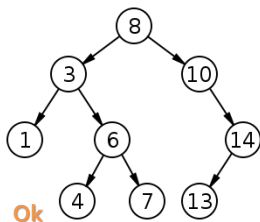


Possible Binary (Search) Trees

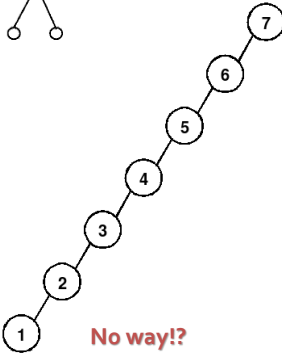
Yes!



Ok



No way!?




Joseph Vella - Practical DBMS

7

Tree Indexing - B+ Tree

7



Other Issues with ISAM Indices:

A *strict* structure and data placement leads to overflow

Index
Fixed area

Overflow pages

Ad hoc growth area: placement is serial

Joseph Vella - Practical DBMS

8

Tree Indexing - B+ Tree

8




Requirements (or desiderata) and Issues

- Tree structure is to be balanced;
 - And *bushy*;
- Tree structure has all leafs at the same level;
- Block based;
- Allows direct and sequential reads;
- All activities have a similar cost basis;
- High space utilisation (per block);
- Resilient to some abuse;
- Sharing mechanism is possible and reasonable.

Joseph Vella - Practical DBMS 9 Tree Indexing - B+ Tree

9




B-Tree Motivation and Approach

- The key technique required to move away from binary search trees (and their variants) to a new solution, B-trees, in that we can choose to *build trees upwards from the bottom instead of downwards from the top*.
 - We have seen how binary trees are stuck with the root node!
- **Bayer and McCreight 1972** (at that time employed by Boeing's but "B" is for Bayer the supervisor) recognized that rather than keeping a bad root and having ever more complex re-arrangement algorithms one should avoid having the root node "anchored down" at the first place!
 - *The B-trees allows the better root to emerge from the key's present while keeping the structure balanced.*

Joseph Vella - Practical DBMS 10 Tree Indexing - B+ Tree

10




B-Tree Characteristics

- **B-trees genera of data structures lend themselves perfectly to paging!**
 - Performance is further enhanced when the root page (and the higher levels perhaps) is kept in ultra fast storage memory (for example in a RAM cache or very fast hard disk).
- **B-trees computational requirements are more on the deterministic side:**
 - can work out the index size and number of levels by knowing the number of keys and page size (and hence order);
 - can work out the range of min and max “time” required to compute operations over B-tree index data structure.
 - **Caution: some of the algorithms look complicated, and indeed they are, but their computational costs is not related to algorithms’ structure.**
- **B-trees have been in use since 1972.**
Informatics have not been the same since then!?

Joseph Vella - Practical DBMS
11
Tree Indexing - B+ Tree

11




B-Tree: Basic Definitions

- **Basic definitions:**
 - **order**
 - **is the maximum number of descendants a node/page can have.**
 - for example, if a page can hold at most seven keys and eight descendants and a minimum of three keys and four descendent then its *order* is eight.
 - **root**
 - **It’s the first page to read and where the first mutli-way split – the higher the order the more partitions possible.**
 - **Assuming that the root’s key are *good* partitions.**
 - **leaves**
 - **are the pages at the lowest level of the B-tree; they do not have any descendants.**

Joseph Vella - Practical DBMS
12
Tree Indexing - B+ Tree

12



B-Tree: Simple example (order = 3)

Search keys are unique

(b)

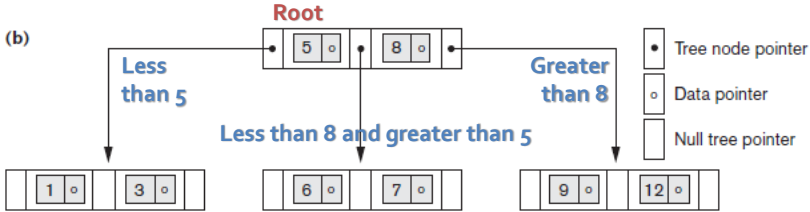



Figure 18.10
B-tree structures. (a) A node in a B-tree with $q - 1$ search values. (b) A B-tree of order $p = 3$. The values were inserted in the order 8, 5, 1, 7, 3, 12, 9, 6.

Joseph Vella - Practical DBMS

13

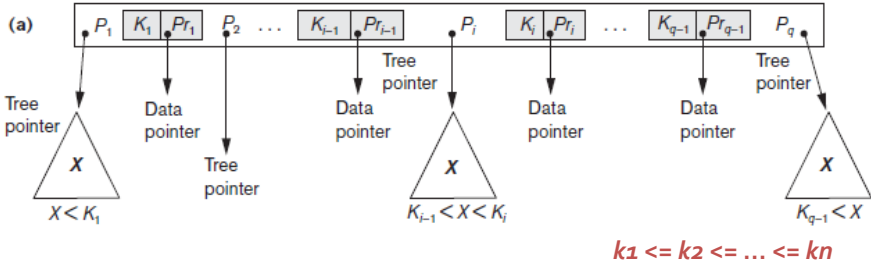
Tree Indexing - B+ Tree

13



B-Tree: Page Structure

(a)



$k_1 \leq k_2 \leq \dots \leq k_n$

Data Structure

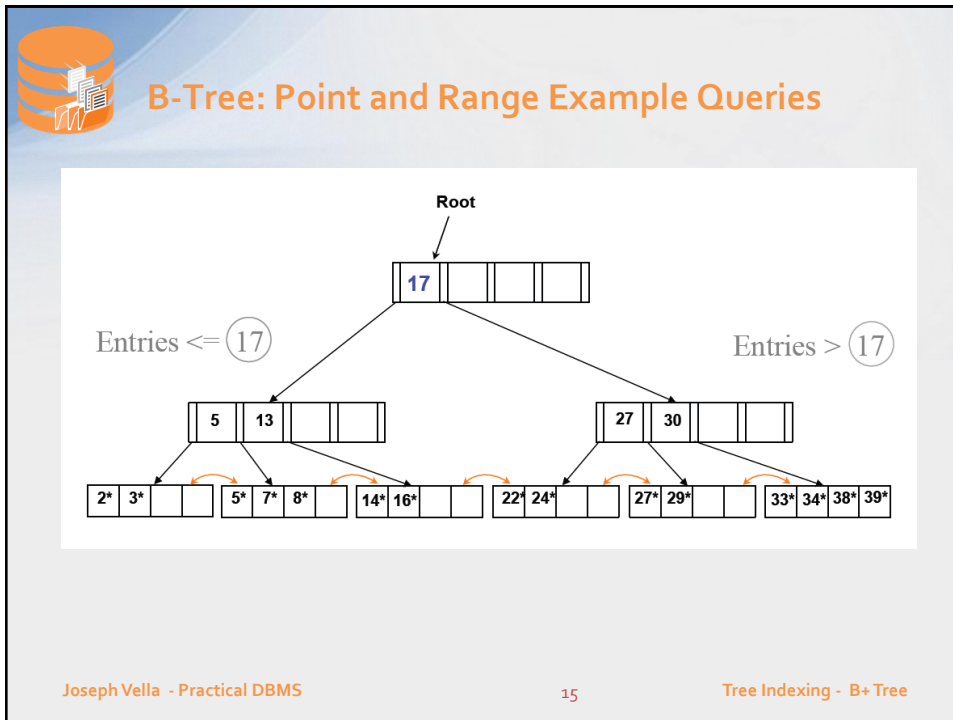
```
TYPE
  BTPAGE      =
  RECORD
    KEYCOUNT : integer;
    KEY        : array[1.. MAXKEYS] of char;
    CHILD      : array[1.. MAXCHILDREN] of integer;
  END;
```

Joseph Vella - Practical DBMS

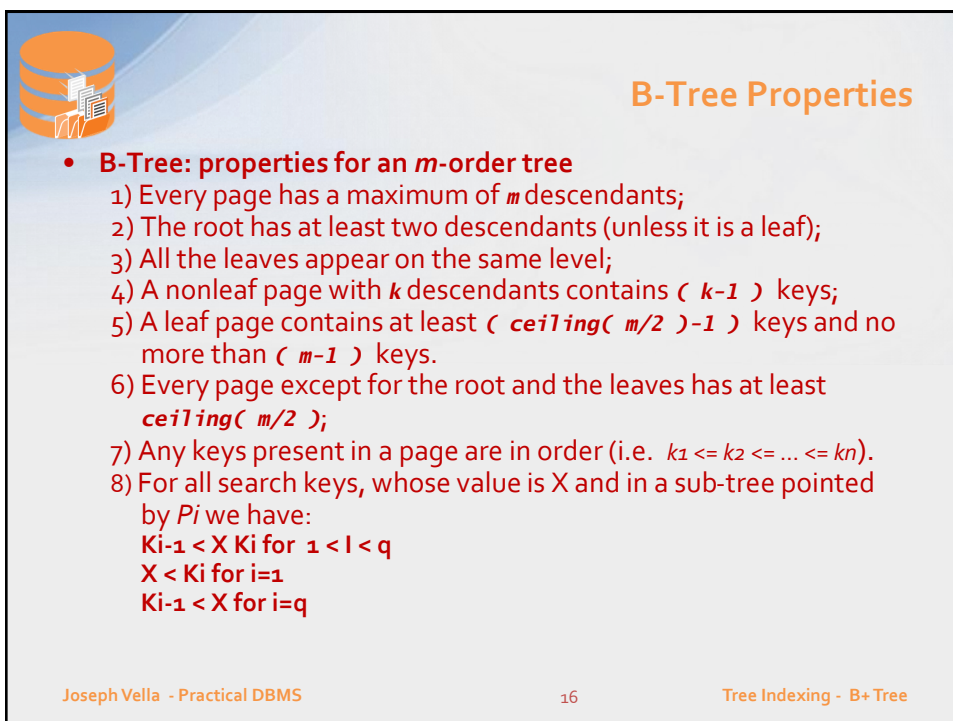
14

Tree Indexing - B+ Tree


14



15



16



B-Trees Deliver


- B-Trees are balanced;
- B-Trees are shallow (distance from root node to any leaf page is short) and consequently require less disk reads;
- Random insertions and deletions in B-Trees are accommodated at a reasonable and predictable costs;
- At least loads 50% use of claimed space - empirical tests put the loading average value at 66%;
- Sequential searches (across pages) not implicit but possible with some effort!
 - look up some details on B+.
- B-Trees are flexible to – can accommodate non-key search field values
 - we need to change the definition of pointer Pri.

Joseph Vella - Practical DBMS

17

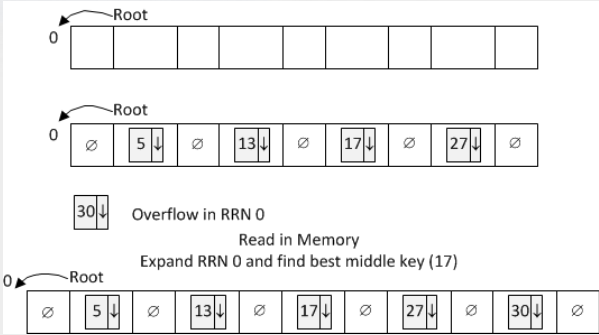
Tree Indexing - B+ Tree

17



B-Tree: Basic Operation – Split & Promote

- B-Tree with order set to 5 (therefore 4 index entries)
- Assume the following key order:
 - 27 [RRNi] , 17 [...], 13 [...], 5 [...], 30 [...]




Joseph Vella - Practical DBMS

18

Tree Indexing - B+ Tree

18



B-Tree: Basic Operation – Split & Promote

Create another two index pages (1,2)

Root

0

∅	5↓	∅	13↓	∅	17↓	∅	27↓	∅	30↓	∅
---	----	---	-----	---	-----	---	-----	---	-----	---

1

--	--	--	--	--	--	--	--	--	--	--

2

--	--	--	--	--	--	--	--	--	--	--

Move later half of page 0 to page 1, and middle key to page 2

Root

0

∅	5↓	∅	13↓	∅	17↓					
---	----	---	-----	---	-----	--	--	--	--	--

1

∅	27↓	∅	30↓	∅						
---	-----	---	-----	---	--	--	--	--	--	--

2


∅	17↓	∅								
---	-----	---	--	--	--	--	--	--	--	--

Joseph Vella - Practical DBMS

19

Tree Indexing - B+ Tree

19



B-Tree: Basic Operation – Split & Promote

Set pointers in page 2 and make page root the new root.

Root

0

∅	5↓	∅	13↓	∅						
---	----	---	-----	---	--	--	--	--	--	--

1

∅	27↓	∅	30↓	∅						
---	-----	---	-----	---	--	--	--	--	--	--

2

•	17↓	•								
---	-----	---	--	--	--	--	--	--	--	--


Write to disk

Joseph Vella - Practical DBMS

20

Tree Indexing - B+ Tree

20




Functions over a B-Tree - Searching

- **Searching**
 - The efficiency being due to the fact that each page is a multi way discriminator (the order m at best).
 - Could be implemented by reading a block into a single main memory buffer.
- **See attached PDF!**
 - For examples & code (pseudo).

Joseph Vella - Practical DBMS 21 Tree Indexing - B+ Tree

21




B Tree – Operation Insert

- **Insertion & Splitting**
 - Is a relatively involved procedure; especially when a page becomes saturated and needs to be split into two pages with the “medium” key promoted. The promoted key might entail, recursively, other splits.
 - A small number of disk pages (for example three) need to be read concurrently into main memory and at worst there can be as many writes as the depth of the tree.
- **See attached PDF!**
 - For examples & code (pseudo).

Joseph Vella - Practical DBMS 22 Tree Indexing - B+ Tree

22




B-Tree Operation - Delete

- **Deletion**
 - Definitely the trickiest part of the implementation.
 - If a non leaf key needs to be purged, the next key (by logical order) has to raise and fill the gap - thus a balanced organisation is maintained.
 - Also, if a page's key count is below the acceptable level then re-grouping to the preceding (by logical order) page must be initialised.
 - A small number of pages have to be read concurrently into main memory. The number of writes equal that of inserts.
- **See attached PDF!**
 - For examples & code (pseudo).

Joseph Vella - Practical DBMS 23 Tree Indexing - B+ Tree

23




B-Tree – Cost of Operations

- **All of these operation require time proportional to the depth of the B tree.**
 - That is $O(\log_n m)$.

Joseph Vella - Practical DBMS 24 Tree Indexing - B+ Tree

24




Working out the worst-case search depth

- Given a B-tree, of order m , and a large number of keys, N , what is the worst-case scenario as regards the tree's depth, d ?
 - For example, say we have a million keys and a 512 ordered key; what will be the max number of page reads required to find a key's entry in the index.
- Observation One:**
 - The number of descendants a certain tree level has is equal to the number of keys present in that level and all the other higher levels, plus one.
 - see attached figure!
- Observation Two:**
 - What is the min number of descendants that any level can have?
 - This permutation yields a B tree that has the max height and the min of breath.
 - Specifically:
 - the root page has 2 as the min;
 - the second level has $\text{ceiling}(m/2)$ per page (i.e., 2);
 - ...
 - the n th level has $\text{ceiling}(m/2)^{d-1} * 2$.
- Workings:**
 - $N+1 \geq \text{ceiling}(m/2)^{d-1} * 2$
 - solving for d ;
 - $d \leq 1 + \log_{\text{ceiling}(m/2)} ((N+1)/2)$
 - $d \leq 1 + \log_{256} 500000.5$
 - $d \leq 3.366$

Joseph Vella - Practical DBMS

25

Tree Indexing - B+ Tree



B Trees vs B+ Trees

B Tree


B+ Tree

<https://www.differencebetween.info/difference-between-b-tree-and-b-plus-tree>

Joseph Vella - Practical DBMS

26

Tree Indexing - B+ Tree




PostgreSQL and B+ indices

- **To create an index in PostgreSQL:**
`CREATE TABLE emp(...);`
 - Note that an index is built on any PK defined.`CREATE INDEX idx_emp_job ON emp(job);``CREATE INDEX idx_emp_deptno_job ON emp(deptno,job);`

Joseph Vella - Practical DBMS 27 Tree Indexing - B+ Tree

27




Index Selection Issues

- **Consider emp table (especially attributes empno, sal, job and deptno)**
 - Intense usage of:
`SELECT *`
`FROM emp`
`WHERE job = $bind_variable;`
 - Occasional usage of:
`SELECT *`
`FROM emp`
`WHERE deptno=$bind_variable;`
 - **Which indexes earn their keep?**
Surely index on job.
If only it's the query on deptno but there is plenty of data in emp then index on deptno (and job) seem to be useful

Joseph Vella - Practical DBMS 28 Tree Indexing - B+ Tree

28




Index Selection Issues

- **Heavy use of the following comparison query:**
 - `SELECT *`
`FROM emp`
`WHERE deptno > $bind_var1`
`AND deptno <= $bind_var2;`
 - `INSERT INTO emp`
`VALUES (...);`
 - Index on deptno, job would be very useful (consider the prefix only).
 - Every index need updating.
- **Heavy use of:**
 - `SELECT *`
`FROM emp`
`WHERE deptno = $bind_var1`
`AND job = $bind_var2;`
 - **Yes ...**

Joseph Vella - Practical DBMS
29
Tree Indexing - B+ Tree

29





B-Tree and batch/bulk uploads

- **Case study:**
 - Is it possible?

Joseph Vella - Practical DBMS
30
Tree Indexing - B+ Tree

30





Joseph Vella - Practical DBMS

31

Tree Indexing - B+ Tree